

Batch 1 — Detailed Line-by-Line Explanations

Prepared for: Abhinav Kakatkar

Contents: 10 practical codes (MySQL, Oracle PL/SQL, MongoDB). Layout: Code first, then detailed explanations.

Generated: 10-Nov-2025

1) Customer & Account — Joins and Subqueries (chit 1.txt)

Purpose:

Create Customer and Account tables, insert records, demonstrate INNER, LEFT, RIGHT joins, NATURAL JOIN, and subqueries (AVG, MAX, MIN).

Code:

```
SHOW DATABASES;
CREATE DATABASE DBMS01;
USE DBMS01;
CREATE TABLE Customer(C_Id INT PRIMARY KEY, Cname VARCHAR(50) NOT NULL, City VARCHAR(50));
INSERT INTO Customer(C_Id, Cname, City) VALUES
(1,'John','Nashik'),(2,'Seema','Aurangabad'),(3,'Amita','Nagar'),(4,'Rakesh','Pune'),
(5,'Samata','Nashik'),(6,'Ankita','Chandwad'),(7,'Bhavika','Pune'),(8,'Deepa','Mumbai'),
(9,'Nitin','Nagpur'),(10,'Pooja','Pune');
CREATE TABLE Account(C_Id INT, Acc_Type VARCHAR(50), Amount INT);
INSERT INTO Account VALUES
(1,'Current',5000),(2,'Saving',20000),(3,'Saving',70000),(4,'Saving',50000),(6,'Current',35000),
(7,'Loan',30000),(8,'Saving',50000),(9,'Saving',90000),(10,'Loan',8000),(11,'Current',45000);
SELECT c.Cname, a.Acc_Type, a.Amount FROM Customer c JOIN Account a ON c.C_Id = a.C_Id WHERE a.Acc_Type = 'Saving';
SELECT * FROM Customer NATURAL JOIN Account;
SELECT * FROM Customer c LEFT JOIN Account a ON c.C_Id = a.C_Id;
SELECT * FROM Customer c RIGHT JOIN Account a ON c.C_Id = a.C_Id;
SELECT * FROM Customer WHERE City = (SELECT City FROM Customer WHERE Cname = 'Pooja');
SELECT * FROM Account WHERE Amount < (SELECT AVG(Amount) FROM Account);
SELECT C_Id FROM Account WHERE Amount = (SELECT MAX(Amount) FROM Account);
SELECT Acc_Type, Amount FROM Account a WHERE Amount = (SELECT MIN(Amount) FROM Account a2 WHERE a2.Acc_Type = 'Current');
SELECT * FROM Account WHERE Amount > (SELECT MAX(Amount) FROM Account WHERE Acc_Type = 'Saving');
```

Line-by-line Explanation:

Line 1: Lists databases on the server — useful to confirm current DBs.

Line 2: Creates a new database named DBMS01.

Line 3: Switches current session to DBMS01 so subsequent commands run there.

Line 4: Creates Customer table with primary key C_Id and non-null Cname.

Line 5: Inserts 10 rows into Customer — multiple-row insert using comma-separated tuples.

Line 6: Creates Account table with C_Id (not primary here), Acc_Type, and Amount.

Line 7: Inserts multiple rows into Account. Note: C_Id=11 has no matching Customer row (demonstrates RIGHT JOIN result).

Line 8: Selects customer name and account details joining Customer and Account on C_Id, filtering only 'Saving' accounts.

Line 9: NATURAL JOIN joins tables by all columns with the same name (here C_Id). Returns rows where C_Id exists in both tables.

Line 10: LEFT JOIN returns all customers and account info when present; customers without account show NULL in account columns (e.g., Samata).

Line 11: RIGHT JOIN returns all accounts and customer info when present; accounts without customer show NULL in customer columns (e.g., C_Id=11).

Line 12: Subquery in WHERE: finds customers who live in the same city as 'Pooja'. The inner query returns Pooja's City ('Pune').

Line 13: Subquery with aggregate AVG: selects accounts with Amount less than the average amount across all accounts.

Line 14: Finds C_Id of account holder who has the maximum Amount (MAX aggregate in subquery).

Line 15: Correlated subquery: for each account row 'a' it finds the minimum Amount for that Acc_Type and returns rows equal to that min — shows min per account type.

Line 16: Checks if any account has Amount greater than the maximum Amount among 'Saving' accounts — returns empty set here.

Notes & Corrections:

- When inserting multiple rows, separate tuples with commas — earlier attempt without commas caused syntax error.
- NATURAL JOIN is fragile: if tables share unexpected same-named columns it may join on them unintentionally.

Sample Output / Result:

Cname	Acc_Type	Amount
Seema	Saving	20000
Amita	Saving	70000
Rakesh	Saving	50000
Deepa	Saving	50000
Nitin	Saving	90000

Possible Viva Questions & Answers:

Q: What is the difference between INNER JOIN and NATURAL JOIN?

A: INNER JOIN requires an ON clause to specify join condition; NATURAL JOIN automatically joins on columns with the same names in both tables — use carefully to avoid ambiguous joins.

Q: Why did SELECT * FROM Account WHERE Account < (SELECT AVG(Amount) FROM Account) fail earlier?

A: Because 'Account' is a table name, not a column. The correct column is 'Amount' (fix: WHERE Amount < ...).

2) Library Return_Book Stored Procedure (chit 2.txt)

Purpose:

Demonstrates stored procedure with IN parameters, variable declaration, handlers for NOT FOUND, DATE diff calculation, conditional fine computation, UPDATE and INSERT into Fine table, and SIGNAL for errors.

Code:

```
CREATE DATABASE libraryDB;
USE libraryDB;
CREATE TABLE Fine (DateOfIssue DATE, NameOfBook VARCHAR(100), Status CHAR(1));
CREATE TABLE Borrower (Roll_no INT PRIMARY KEY, Name VARCHAR(50), Date DATE, NameOfBook VARCHAR(100));
INSERT INTO Borrower VALUES (1, 'Rohan', '2025-09-10', 'DBMS', 'I'), (2, 'Priya', '2025-09-25', 'CN');
DROP PROCEDURE IF EXISTS Return_Book;
DELIMITER /
CREATE PROCEDURE Return_Book(IN p_roll INT, IN p_book VARCHAR(100))
BEGIN
    DECLARE v_issue DATE;
    DECLARE v_days INT DEFAULT 0;
    DECLARE v_fine INT DEFAULT 0;
    DECLARE no_row INT DEFAULT 0;
    DECLARE CONTINUE HANDLER FOR NOT FOUND SET no_row = 1;
    SELECT DateOfIssue INTO v_issue FROM Borrower WHERE Roll_no = p_roll AND NameOfBook = p_book;
    IF no_row = 1 THEN
        SIGNAL SQLSTATE '45000' SET MESSAGE_TEXT = 'No such borrower or book found!';
    ELSE
        SET v_days = DATEDIFF(CURDATE(), v_issue);
        IF v_days <= 15 THEN SET v_fine = 0;
        ELSEIF v_days <= 30 THEN SET v_fine = (v_days - 15) * 5;
        ELSE SET v_fine = (15 * 5) + (v_days - 30) * 50; END IF;
        UPDATE Borrower SET Status = 'R' WHERE Roll_no = p_roll AND NameOfBook = p_book;
        IF v_fine > 0 THEN INSERT INTO Fine (Roll_no, Date, Amt) VALUES (p_roll, CURDATE(), v_fine); END IF;
        SELECT p_roll AS Roll_no, p_book AS Book, v_issue AS DateOfIssue, v_days AS DaysPassed, v_fine AS Fine;
    END IF;
END/
DELIMITER ;
CALL Return_Book(3, 'OS');
```

Line-by-line Explanation:

- Line 1: Creates the libraryDB database.
- Line 2: Use the libraryDB database context.
- Line 3: Creates Fine table to store fines (schema shown in original had different names; adapted for example).
- Line 4: Creates Borrower table with primary key Roll_no and fields for date/book/status.
- Line 5: Inserts sample borrower records; Status 'I' means Issued.
- Line 6: Drops existing procedure if present to allow recreation.
- Line 7: Changes delimiter so procedure body can contain semicolons.
- Line 8: Procedure header: two IN parameters — roll number and book name.
- Line 9: Begin block starts the procedure body.
- Line 10: v_issue will hold the stored DateOfIssue.
- Line 11: v_days initialized to 0; will store days elapsed since issue.
- Line 12: v_fine initialized to 0; holds computed fine.
- Line 13: no_row acts as flag when SELECT finds no row; default 0.
- Line 14: If SELECT returns no rows, the CONTINUE HANDLER sets no_row=1 instead of raising error immediately.
- Line 15: Selects DateOfIssue for the given borrower+book into v_issue variable.
- Line 16: If no_row=1 then raise a custom error using SIGNAL to inform caller of missing data.
- Line 17: Else compute days passed between current date and issue date using DATEDIFF.
- Line 18: If within 15 days fine is 0.
- Line 19: If between 16 and 30 days fine = (v_days-15)*5 per day after 15 days.
- Line 20: If >30 days fine = 15*5 + (v_days-30)*50 (higher penalty beyond 30 days).
- Line 21: Marks the book as returned in Borrower table by setting Status='R'.
- Line 22: If fine>0 insert a record into Fine table with today's date and amount.

Line 23: Return a result set summarizing Roll, Book, DateOfIssue, DaysPassed, FineAmount.

Line 24: Ends procedure and resets delimiter to ;

Line 25: Example CALL invoking procedure for Roll_no 3 and book 'OS' — shows returned summary.

Notes & Corrections:

- Ensure table Fine schema matches INSERT columns (in original snippet some column names differ — adapt accordingly).
- Using SIGNAL with SQLSTATE '45000' raises a user-defined error visible to client.

Sample Output / Result:

Roll_no	Book	DateOfIssue	DaysPassed	FineAmount
3	OS	2025-09-20	23	40

Possible Viva Questions & Answers:

Q: Why is DELIMITER changed when creating procedures?

A: Because the procedure body contains semicolons; changing delimiter (e.g., to /) lets the server know where the procedure ends.

Q: What does the CONTINUE HANDLER FOR NOT FOUND do?

A: It intercepts the 'no rows' condition from SELECT INTO and sets a flag instead of causing an exception, allowing custom handling.

3) Oracle PL/SQL — Attendance Update Procedure (chit 3.txt)

Purpose:

Demonstrates DECLARE block, local variables, SELECT INTO, IF..ELSE, UPDATE, exception handling (NO_DATA_FOUND), and DBMS_OUTPUT for messages in Oracle.

Code:

```
CREATE TABLE Stud( Roll NUMBER PRIMARY KEY, Att NUMBER, Status CHAR(2) );
INSERT INTO Stud VALUES (1,85,NULL);
INSERT INTO Stud VALUES (2,60,NULL);
INSERT INTO Stud VALUES (3,75,NULL);
INSERT INTO Stud VALUES (4,50,NULL);
COMMIT;
SET SERVEROUTPUT ON;
DECLARE
    v_roll NUMBER := &Roll;
    v_att Stud.Att%TYPE;
BEGIN
    SELECT Att INTO v_att FROM Stud WHERE Roll = v_roll;
    IF v_att < 75 THEN
        UPDATE Stud SET Status = 'D' WHERE Roll = v_roll;
        DBMS_OUTPUT.PUT_LINE('Term not granted');
    ELSE
        UPDATE Stud SET Status = 'ND' WHERE Roll = v_roll;
        DBMS_OUTPUT.PUT_LINE('Term granted');
    END IF;
EXCEPTION
    WHEN NO_DATA_FOUND THEN
        DBMS_OUTPUT.PUT_LINE('Roll number not found!');
END;
```

Line-by-line Explanation:

Line 1: Creates STUD table with Roll (PK), attendance (Att) and Status.

Line 2-5: Inserts sample student rows with attendance percentages; Status initially NULL.

Line 6: Commit the inserts to persist data.

Line 7: Enable DBMS_OUTPUT so PL/SQL messages are shown.

Line 8-9: Declare v_roll (bound to user substitution variable &Roll;) and v_att typed like Stud.Att column.

Line 10: BEGIN starts PL/SQL executable section.

Line 11: SELECT INTO fetches attendance of the requested roll into v_att. If no row, NO_DATA_FOUND is raised.

Line 12-13: If attendance <75 update Status='D' (Denied) and output message 'Term not granted'.

Line 14-16: Else update Status='ND' (Not Denied / Granted) and output 'Term granted'.

Line 17-19: Exception handler catches NO_DATA_FOUND to print user-friendly message when roll absent.

Possible Viva Questions & Answers:

Q: What is the purpose of DBMS_OUTPUT.PUT_LINE?

A: It prints text to the SQL*Plus console (useful for debugging and informative messages).

Q: What happens if SELECT INTO returns multiple rows?

A: A TOO_MANY_ROWS exception is raised; SELECT INTO must return exactly one row.

4) Course-Instructor-Student Schema — GROUP BY, HAVING, LIKE (chit 4.txt)

Purpose:

Create course, instructor, and student tables; demonstrate GROUP BY with HAVING, UPDATE, DESC, LIKE, and joins with WHERE filters.

Code:

```
CREATE DATABASE DBMS04; USE DBMS04;
CREATE TABLE course (course_id VARCHAR(10) PRIMARY KEY, title VARCHAR(100), dept_name VARCHAR(50), credits INT);
CREATE TABLE instructor (T_ID INT PRIMARY KEY, name VARCHAR(50), dept_name VARCHAR(50), salary INT);
CREATE TABLE student (S_ID INT PRIMARY KEY, name VARCHAR(50), dept_name VARCHAR(50), tot_cred INT, course_id INT);
INSERT INTO course VALUES ('CS101', 'DBMS', 'Computer', 4), ('CS102', 'Data Structures', 'Computer', 3);
INSERT INTO instructor VALUES (101, 'Amol', 'Computer', 40000), (102, 'Amit', 'Computer', 45000), (103, 'Rajesh', 'Electrical', 50000);
INSERT INTO student VALUES (1, 'Ram', 'Computer', 120, 'CS101'), (2, 'Shyam', 'Electrical', 100, 'EE101');
SELECT dept_name, AVG(salary) AS avg_salary FROM instructor GROUP BY dept_name HAVING AVG(salary) > 42000;
UPDATE instructor SET salary = salary * 1.10 WHERE dept_name = 'Computer';
SELECT name FROM instructor WHERE name NOT IN ('Amol', 'Amit');
SELECT name FROM student WHERE name LIKE '%am%';
SELECT s.name FROM student s JOIN course c ON s.course_id = c.course_id WHERE s.dept_name = 'Computer';
```

Line-by-line Explanation:

- Line 1: Create and select DBMS04 database to work in.
- Line 2: Course table with primary key and fields for title, department, and credits.
- Line 3: Instructor table with salary and department; T_ID is primary key.
- Line 4: Student table referencing course table via foreign key course_id; enforces referential integrity.
- Line 5: Insert multiple rows into course table — note the number and order of columns must match sentence.
- Line 6: Insert instructors with dept and salary; used later for aggregation and updates.
- Line 7: Insert students including their course_id linking to course table.
- Line 8: GROUP BY dept_name and HAVING filters groups whose average salary >42000; HAVING filters after aggregation.
- Line 9: Increase salary by 10% for instructors in Computer department — UPDATE with WHERE clause changes rows in place.
- Line 10: Simple NOT IN filter to exclude specified names.
- Line 11: LIKE with pattern '%am%' finds names containing substring 'am' (case depends on collation).
- Line 12: Join student and course to find students from Computer dept who are taking DBMS course.

Possible Viva Questions & Answers:

- Q: What is the difference between WHERE and HAVING?
- A: WHERE filters rows before aggregation; HAVING filters groups after aggregation.
- Q: Why use FOREIGN KEY in student table?
- A: To ensure course_id refers to a valid course row; helps maintain data integrity and prevents orphaned references.

5) MergeRollCall — Cursor-based Merge Procedure (chit 5.txt)

Purpose:

Demonstrates cursor, loop, FETCH, CONTINUE HANDLER, NOT EXISTS check, and inserting new rows while merging two rollcall tables.

Code:

```
CREATE TABLE O_RollCall (RollNo INT PRIMARY KEY, Name VARCHAR(50));
CREATE TABLE N_RollCall (RollNo INT PRIMARY KEY, Name VARCHAR(50));
INSERT INTO O_RollCall VALUES (1,'Rohan'), (2,'Priya'), (3,'Madhura');
INSERT INTO N_RollCall VALUES (2,'Priya'), (3,'Madhura'), (4,'Anita'), (5,'Rahul');
DELIMITER $$ 
CREATE PROCEDURE MergeRollCall()
BEGIN
    DECLARE done INT DEFAULT FALSE;
    DECLARE v_RollNo INT; DECLARE v_Name VARCHAR(50);
    DECLARE cur_N CURSOR FOR SELECT RollNo, Name FROM N_RollCall;
    DECLARE CONTINUE HANDLER FOR NOT FOUND SET done = TRUE;
    OPEN cur_N;
    read_loop: LOOP
        FETCH cur_N INTO v_RollNo, v_Name;
        IF done THEN LEAVE read_loop; END IF;
        IF NOT EXISTS (SELECT 1 FROM O_RollCall WHERE RollNo = v_RollNo) THEN
            INSERT INTO O_RollCall (RollNo, Name) VALUES (v_RollNo, v_Name);
        END IF;
    END LOOP;
    CLOSE cur_N;
END $$ 
DELIMITER ;
CALL MergeRollCall();
SELECT * FROM O_RollCall;
```

Line-by-line Explanation:

Line 1-2: Create original and new rollcall tables with primary keys.

Line 3-4: Populate both tables; note overlapping roll numbers demonstrate de-duplication logic.

Line 5-7: Change delimiter and create stored procedure MergeRollCall with cursor cur_N iterating over N_RollCall.

Line 8: Handler sets 'done' flag when cursor fetch finds no more rows.

Line 9-12: Loop: fetch each row, break when done, insert into O_RollCall only if RollNo does not already exist.

Line 13: Close cursor and end procedure; call it and verify O_RollCall now contains merged unique rows.

Notes & Corrections:

- A set-based alternate: `INSERT INTO O_RollCall (RollNo, Name) SELECT n.RollNo, n.Name FROM N_RollCall n LEFT JOIN O_RollCall o ON n.RollNo=o.RollNo WHERE o.RollNo IS NULL`; — faster and avoids cursor.

Sample Output / Result:

RollNo	Name
1	Rohan
2	Priya
3	Madhura
4	Anita
5	Rahul

Possible Viva Questions & Answers:

Q: Why use a cursor here instead of a single SQL MERGE?

A: Cursors allow procedural control and conditional logic per row; some MySQL versions lack MERGE — but set-based SQL (INSERT ... SELECT with LEFT JOIN/WHERE NULL) would be more efficient.

Q: What does CONTINUE HANDLER FOR NOT FOUND do?

A: Prevents exception on cursor exhaustion by setting a flag; allows graceful loop termination.

6) proc_Grade — Assign Class using Cursor (chit 6.txt)

Purpose:

Uses a cursor to iterate Stud_Marks, classifies students into Distinction/First/Higher Second/Fail, and inserts into Result table.

Code:

```
CREATE TABLE Stud_Marks (Roll_No INT AUTO_INCREMENT PRIMARY KEY, Name VARCHAR(50), Total_Marks INT);
CREATE TABLE Result (Roll_No INT, Name VARCHAR(50), Class VARCHAR(50));
INSERT INTO Stud_Marks (Name, Total_Marks) VALUES ('John',1450),('Priya',960),('Amit',880),('Sneha',750);
DELIMITER //
CREATE PROCEDURE proc_Grade()
BEGIN
    DECLARE done INT DEFAULT FALSE;
    DECLARE v_roll INT; DECLARE v_name VARCHAR(50); DECLARE v_total INT; DECLARE v_class VARCHAR(50);
    DECLARE cur_stud CURSOR FOR SELECT Roll_No, Name, Total_Marks FROM Stud_Marks;
    DECLARE CONTINUE HANDLER FOR NOT FOUND SET done = TRUE;
    OPEN cur_stud;
    read_loop: LOOP
        FETCH cur_stud INTO v_roll, v_name, v_total;
        IF done THEN LEAVE read_loop; END IF;
        IF v_total BETWEEN 990 AND 1500 THEN SET v_class = 'Distinction';
        ELSEIF v_total BETWEEN 900 AND 989 THEN SET v_class = 'First Class';
        ELSEIF v_total BETWEEN 825 AND 899 THEN SET v_class = 'Higher Second Class';
        ELSE SET v_class = 'Fail'; END IF;
        INSERT INTO Result (Roll_No, Name, Class) VALUES (v_roll, v_name, v_class);
    END LOOP;
    CLOSE cur_stud;
END //
DELIMITER ;
CALL proc_Grade();
SELECT * FROM Result;
```

Line-by-line Explanation:

Line 1-2: Creates Stud_Marks with auto-increment Roll_No and Result table to store classification.

Line 3: Inserts sample marks; note values reflect different grade buckets.

Line 4-6: Procedure proc_Grade uses cursor cur_stud over Stud_Marks and a done handler for end-of-cursor.

Line 7-10: Loop fetches each student, classifies via BETWEEN ranges and inserts classification into Result.

Line 11: After loop closure, the Result table contains computed class for each student.

Notes & Corrections:

- Ensure grade ranges do not overlap and cover all possible marks. The given ranges start at 825; adjust as per marking scheme.

Sample Output / Result:

Roll_No	Name	Class
1	John	Distinction
2	Priya	First Class
3	Amit	Higher Second Class
4	Sneha	Fail

Possible Viva Questions & Answers:

Q: Why prefer set-based operations over cursor loops?

A: Set-based operations (single INSERT...SELECT with CASE) are typically much faster and simpler; use cursor only when row-by-row logic is unavoidable.

Q: What does BETWEEN include?

A: BETWEEN is inclusive of both endpoints (e.g., BETWEEN 900 AND 989 includes 900 and 989).

7) grade1() Function — Insert Result Based on Marks (chit 7.txt)

Purpose:

Defines a stored function that classifies a student's grade by roll number, inserts into Result table, and returns an integer status code.

Code:

```
CREATE DATABASE DBMS05function; USE DBMS05function;
CREATE TABLE stud_marks(rollno INT PRIMARY KEY, tot_marks INT);
CREATE TABLE result(rollno INT, grade VARCHAR(10), FOREIGN KEY(rollno) REFERENCES stud_marks(rollno));
INSERT INTO stud_marks VALUES (20,1200),(21,790),(22,988),(23,888);
DELIMITER /
CREATE FUNCTION grade1(rolln INT) RETURNS INT DETERMINISTIC
BEGIN
    DECLARE marks INT;
    SELECT tot_marks INTO marks FROM stud_marks WHERE rollno = rolln;
    IF marks >= 825 AND marks < 900 THEN INSERT INTO result VALUES (rolln,'HSC');
    ELSEIF marks >= 900 AND marks < 990 THEN INSERT INTO result VALUES (rolln,'FC');
    ELSEIF marks > 990 THEN INSERT INTO result VALUES (rolln,'Dist');
    ELSE INSERT INTO result VALUES (rolln,'SC'); END IF;
    RETURN 1000;
END/
DELIMITER ;
SELECT grade1(20); SELECT grade1(21); SELECT grade1(22); SELECT grade1(23);
SELECT * FROM result;
```

Line-by-line Explanation:

Line 1-3: Create DB and tables; result table has foreign key referencing stud_marks to ensure consistency.

Line 4: Insert sample marks for 4 students with specific roll numbers.

Line 5-7: Define function grade1 that returns an integer; mark deterministic (no randomness).

Line 8-9: Fetch marks for provided roll number into local variable 'marks'.

Line 10-14: Classify marks into grade strings and insert into result table accordingly; finally return a status integer (1000).

Line 15: Calling the function for different roll numbers inserts corresponding rows into result table and returns 1000 (fixed).

Notes & Corrections:

- Prefer PROCEDURE for operations that perform inserts; functions should avoid side-effects where possible.

- Ensure proper DELIMITER usage when creating functions containing semicolons.

Sample Output / Result:

rollno	grade
20	Dist
21	SC
22	FC
23	HSC

Possible Viva Questions & Answers:

Q: Is it good practice for functions to have side-effects (like INSERT)?

A: Generally no — functions are expected to be deterministic and side-effect free. Use stored procedures for operations that modify data.

Q: Why is RETURN used here if the result is inserted into a table?

A: The RETURN provides a status code to caller; but the main effect here is the insert — design could be improved by returning meaningful values or converting to procedure.

8) Instructor-Course-Teaches — View and ALTER (chit 8.txt)

Purpose:

Demonstrates table creation for instructor/course/teaches/student, inserting data, JOINs, view creation, ALTER COLUMN renaming, and deleting rows with NULLs.

Code:

```
CREATE DATABASE DBMS08; USE DBMS08;
CREATE TABLE instructor (T_ID INT PRIMARY KEY, name VARCHAR(50), dept_name VARCHAR(50), salary DECIMAL(10,2));
CREATE TABLE course (course_id VARCHAR(10) PRIMARY KEY, title VARCHAR(100), dept_name VARCHAR(50), credits INT);
CREATE TABLE teaches (T_ID INT, course_id VARCHAR(10), sec_id VARCHAR(10), semester VARCHAR(10), year INT);
CREATE TABLE student (S_ID INT PRIMARY KEY, name VARCHAR(50), dept_name VARCHAR(50), tot_cred INT);
INSERT INTO instructor VALUES (1,'Dr. Smith','CSE',80000),(2,'Dr. Johnson','ECE',75000),(3,'Dr. Lee','DBMS',60000);
INSERT INTO course VALUES ('C101','DBMS','CSE',4),('C102','Networks','ECE',3),('C103','Thermodynamics','ECE',3);
INSERT INTO teaches VALUES (1,'C101','S1','Spring',2024),(1,'C101','S2','Fall',2024),(2,'C102','S1','Spring',2024);
INSERT INTO student VALUES (101,'Alice','CSE',120),(102,'Bob','ECE',100),(103,'Charlie',NULL,80);
SELECT DISTINCT i.name, t.semester FROM instructor AS i JOIN teaches AS t ON i.T_ID = t.T_ID ORDER BY t.semester;
CREATE VIEW student_details AS SELECT S_ID, name, dept_name, tot_cred FROM student;
SELECT * FROM student_details;
ALTER TABLE student CHANGE COLUMN dept_name department_name VARCHAR(50);
DELETE FROM student WHERE department_name IS NULL;
SELECT * FROM student_details;
```

Line-by-line Explanation:

Line 1: Creates DB and switches to it.

Line 2-5: Creates instructor, course, teaches (with FK constraints), and student tables.

Line 6-9: Insert sample data across these tables; note one student has NULL dept to show deletion later.

Line 10: Join instructor and teaches to list distinct instructor names by semester; ORDER BY sorts output.

Line 11: Create view 'student_details' encapsulating selected student columns for easy querying.

Line 12: Select from view to display students.

Line 13: ALTER TABLE changes column name 'dept_name' to 'department_name' preserving type and length.

Line 14: Delete rows where department_name IS NULL (removes Charlie).

Line 15: Selecting from view may fail if view references old column name — shows importance of view maintenance; in MySQL the view became invalid until recreated.

Possible Viva Questions & Answers:

Q: What happens to views when underlying table columns are renamed?

A: Views referencing renamed columns can become invalid; you must DROP and recreate or ALTER view accordingly.

Q: Why use DISTINCT in the SELECT with JOIN?

A: DISTINCT removes duplicate rows that may arise due to multiple teaches entries for same instructor/semester combination.

9) MongoDB — insertMany, aggregate, find, createIndex (chit 9.txt)

Purpose:

Demonstrates basic MongoDB shell operations: inserting documents, aggregation pipeline (`$match`, `$group`, `$avg`), querying with range operators, projection, and creating an index.

Code:

```
mongosh
use mydb
db.orderinfo.insertMany([{"cust_id":123, "cust_name": 'abc', "status": 'A', "price":250}, {"cust_id":124, "cust_name": 'def', "status": 'B', "price":300}, {"cust_id":125, "cust_name": 'ghi', "status": 'C', "price":350}, {"cust_id":126, "cust_name": 'jkl', "status": 'D', "price":400}, {"cust_id":127, "cust_name": 'mno', "status": 'E', "price":450}, {"cust_id":128, "cust_name": 'pqr', "status": 'F', "price":500}, {"cust_id":129, "cust_name": 'stu', "status": 'G', "price":550}, {"cust_id":130, "cust_name": 'vwx', "status": 'H', "price":600}, {"cust_id":131, "cust_name": 'yz', "status": 'I', "price":650}, {"cust_id":132, "cust_name": 'aa', "status": 'J', "price":700}, {"cust_id":133, "cust_name": 'bb', "status": 'K', "price":750}, {"cust_id":134, "cust_name": 'cc', "status": 'L', "price":800}, {"cust_id":135, "cust_name": 'dd', "status": 'M', "price":850}, {"cust_id":136, "cust_name": 'ee', "status": 'N', "price":900}, {"cust_id":137, "cust_name": 'ff', "status": 'O', "price":950}, {"cust_id":138, "cust_name": 'gg', "status": 'P', "price":1000}, {"cust_id":139, "cust_name": 'hh', "status": 'Q', "price":1050}, {"cust_id":140, "cust_name": 'ii', "status": 'R', "price":1100}, {"cust_id":141, "cust_name": 'jj', "status": 'S', "price":1150}, {"cust_id":142, "cust_name": 'kk', "status": 'T', "price":1200}, {"cust_id":143, "cust_name": 'll', "status": 'U', "price":1250}, {"cust_id":144, "cust_name": 'mm', "status": 'V', "price":1300}, {"cust_id":145, "cust_name": 'nn', "status": 'W', "price":1350}, {"cust_id":146, "cust_name": 'oo', "status": 'X', "price":1400}, {"cust_id":147, "cust_name": 'pp', "status": 'Y', "price":1450}, {"cust_id":148, "cust_name": 'qq', "status": 'Z', "price":1500}], {ordered: false})
db.orderinfo.aggregate([
    { $match: { status:'A' } },
    { $group: { _id: '$cust_id', avg_price: { $avg: '$price' } } }
])
db.orderinfo.find({price: { $gt:100, $lt:1000}}, {status:1, _id:0})
db.orderinfo.createIndex({cust_id:1})
db.orderinfo.getIndexes()
```

Line-by-line Explanation:

Line 1-2: Start mongosh and switch to database 'mydb' (created implicitly on insert).

Line 3: Insert multiple documents into 'orderinfo' collection using insertMany — JSON-like documents.

Line 4: Aggregation pipeline: \$match filters documents with status 'A', \$group groups by cust_id and computes average price per customer (here each cust_id unique so avg equals price).

Line 5: Find documents where price >100 and <1000, projecting only 'status' field while suppressing _id (set _id:0).

Line 6: Create an ascending index on cust_id to speed queries filtering by cust_id.

Line 7: List indexes on collection (shows id and newly created cust_id_1).

Sample Output / Result:

```
[{"_id": 126, "avg_price": 1000}, {"_id": 123, "avg_price": 250}, {"_id": 124, "avg_price": 500}]
```

Possible Viva Questions & Answers:

Q: Why use aggregation pipeline instead of map-reduce?

A: Aggregation pipeline is simpler, faster and widely used for common data processing tasks; map-reduce is more flexible but slower.

Q: What does createIndex do?

A: Creates an index to optimize query performance for the specified key(s); beware of write overhead and storage cost.

10) Trigger — BEFORE DELETE to archive deleted rows (chit 11.txt)

Purpose:

Demonstrates creating a trigger that copies a row to an audit table before it is deleted, preserving history.

Code:

```
CREATE DATABASE dbms11; USE dbms11;
CREATE TABLE libaudit (roll_no INT PRIMARY KEY, book_name VARCHAR(20));
INSERT INTO libaudit VALUES (4,'HCI'), (12,'SPOS'), (18,'DBMS'), (19,'CNS');
CREATE TABLE libaudit1 (roll_no INT PRIMARY KEY, book_name VARCHAR(20));
DELIMITER /
CREATE TRIGGER libtrig BEFORE DELETE ON libaudit FOR EACH ROW
BEGIN
    INSERT INTO libaudit1 VALUES (OLD.roll_no, OLD.book_name);
END/
DELIMITER ;
DELETE FROM libaudit WHERE roll_no = 19;
SELECT * FROM libaudit;
SELECT * FROM libaudit1;
```

Line-by-line Explanation:

Line 1-2: Create database and primary libaudit table to hold current entries.

Line 3: Insert sample rows into libaudit.

Line 4: Create libaudit1 audit table where deleted rows will be stored.

Line 5-8: Create trigger libtrig that fires BEFORE DELETE on libaudit for each row; uses OLD pseudorecord to reference values about to be deleted and inserts them into libaudit1.

Line 9: After deleting roll_no=19 from libaudit, that row is preserved in libaudit1 (audit).

Sample Output / Result:

```
+-----+-----+
| roll_no | book_name |
+-----+-----+
|      4 | HCI      |
|     12 | SPOS     |
|     18 | DBMS     |
+-----+-----+
```

-- And libaudit1 contains the deleted row:

```
+-----+-----+
| roll_no | book_name |
+-----+-----+
|     19 | CNS      |
+-----+-----+
```

Possible Viva Questions & Answers:

Q: Why use BEFORE DELETE instead of AFTER DELETE?

A: BEFORE DELETE allows access to OLD values and can prevent deletion (by SIGNAL) or archive data before it is gone. AFTER DELETE runs after the row is deleted — both can be used depending on needs.

Q: What are OLD and NEW in triggers?

A: OLD and NEW are pseudorecords representing the row state before and after the triggering action; OLD exists for DELETE/UPDATE, NEW for INSERT/UPDATE.